

## MTH210: Lab 4 Solutions


1. **Implement an AR algorithm to sample from a Beta(2, .1) distribution. Follow the theory from the notes.**

Let us do some theory for this first. The target density is:

$$f(x) = \frac{\Gamma(2.1)}{\Gamma(2)\Gamma(1)} x^1 (1-x)^{1-1} \quad ; 0 \leq x \leq 1.$$

The shape of the target density is

```
x <- seq(0, 1, length = 5e2)
plot(x, dbeta(x, 2, .1), type = 'l', ylab = "Density")
```



As you can see, as  $x \rightarrow 1$ , the target density diverges. Thus, using a uniform proposal will not work in this problem. Instead, we do the same trick as we did in class:

$$f(x) = \frac{\Gamma(2.1)}{\Gamma(2)\Gamma(1)} x^1 (1-x)^{1-1} \Rightarrow g(x) \propto (1-x)^{(.1-1)}.$$

This means,  $g(x)$  is the density of Beta(1, .1). Since the task is to draw from Beta, I do not want to use `rbeta` command. Instead will use inverse transform.

$$g(x) = .1(1-x)^{(.1-1)} \Rightarrow G(x) = \int_0^x g(x)dx = 1 - (1-t)^{(.1)}.$$

So the inverse is

$$G^{-1}(u) = 1 + (1-u)^{1/.1}.$$

Using the above, we can draw our proposal. However, we have to be careful about numerical inaccuracies, so we will do the following:

$$\log(1 - G^{-1}(u)) = 10 \log(1 - u)$$

and then exponentiate.

```
#####
## Accept-reject for
## Beta(2,.1) distribution
## Using Beta(1,.1) proposal
#####
beta_ar <- function(m = 2, n = .1)
{
  c <- gamma(m + n)/(gamma(n)*gamma(m))/n
  accept <- 0
  counter <- 0 # count the number of loop
  while(accept == 0)
  {
    counter <- counter + 1

    # Inverse transform to draw from proposal
    U <- runif(1)
    foo <- (1/n) * log(1 - U)
    prop <- 1 - exp(foo)

    # log ratio
    log.ratio <- dbeta(prop, m, n, log = TRUE) - log(c)
    - dbeta(prop, 1, n, log = TRUE)
    if(log(runif(1)) <= log.ratio)
    {
      accept <- 1
      return(c(prop, counter))
    }
  }
}
```

```

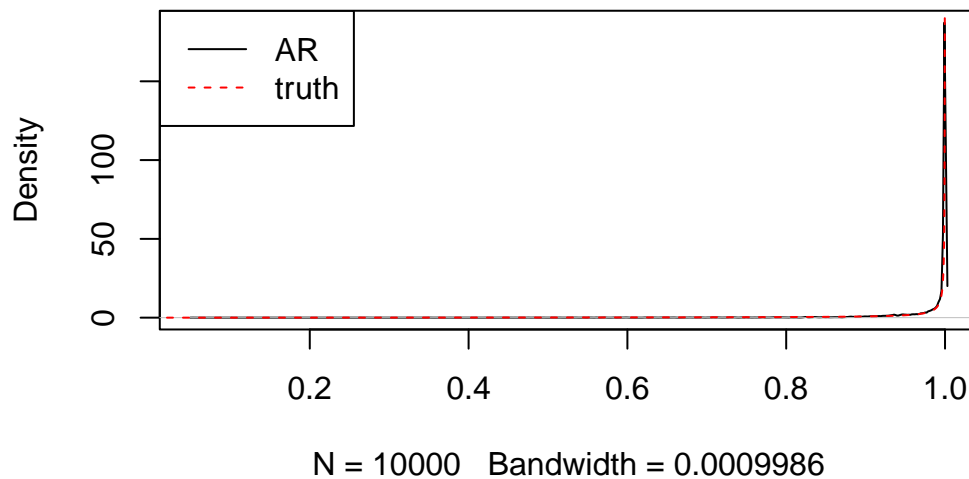
}

### Obtaining 104 samples from Beta() distribution
N <- 1e4
samp <- numeric(length = N)
counts <- numeric(length = N)
for(i in 1:N)
{
  rep <- beta_ar() ## fill in
  samp[i] <- rep[1] ## fill in
  counts[i] <- rep[2] ## fill in
}

# Make a plot of the estimated density from the samples
# versus the true density
x <- seq(0, 1, length = 5000)
plot(density(samp), main = "Estimated density from 1e4 samples")
lines(x, dbeta(x, 2, .1), col = "red", lty = 2) ## Complete this
legend("topleft", lty = 1:2, col = c("black", "red"), legend = c("AR", "truth"))

```

### Estimated density from 1e4 samples



Since the estimated density from my  $10^4$  samples matches the true density, I can be confident that I have coded this correctly. In the above, to avoid numerical instability, instead of calculating

$$\frac{f(y)}{cg(y)}$$

I calculate:

$$\log(f(y)) - \log(c) - \log(g(y))$$

And compare log of the ratio to log of a uniform random variable.

2. Using only  $U(0,1)$  draws, draw samples from  $\text{Gamma}(4,3)$  using Accept-Reject and an exponential proposal. Compare the performance of the sampler using the optimal exponential proposal, versus  $\lambda = 2$ .

We have done this problem theoretically in the class. The way we found the optimal value for an Exponential Proposal was by finding the bound  $c$  as a function of  $c(\lambda)$ , and then determining which  $\lambda$  minimizes  $c(\lambda)$ . First note that  $\alpha > 1$  and so the optimal value of  $\lambda$  in this proposal is

$$\lambda = \frac{\beta}{\alpha} = \frac{4}{3}.$$

Further, from our notes, recall that for any  $\lambda < \beta$ , the  $\sup f(x)/g(x)$  occurs at

$$x = \frac{\alpha - 1}{\beta - \lambda}.$$

```
#####
## Accept-reject for
## Gamma(4,3) distribution
## Using Exp(lambda) proposal
#####
gamma_ar <- function(alpha, beta, lambda)
{
  if(alpha < 1)
  {
    stop("alpha less than 1. AR not possible with Exponential proposal")
  }
  accept <- 0
  counter <- 0 # count the number of loop

  # value of x where f/g is max
  max.x <- (alpha - 1)/(beta - lambda)

  # calculating c at that value +some little thing
  c.lambda <- dgamma(max.x, alpha, beta)/ dexp(max.x, rate = lambda) + .00001
  while(accept == 0)
  {
    counter <- counter + 1

    # from Exp(lambda)
```

```

# using inverse-transform
foo <- runif(1)
prop <- -log(foo)/lambda

log.ratio <- dgamma(prop, alpha, beta, log = TRUE) - dexp(prop, rate = lambda, log = TRUE)

if(log(runif(1)) <= log.ratio)
{
  accept <- 1
  return(c(prop, counter))
}
}
}

```

```

### Obtaining  $10^4$  samples from Beta() distribution
N <- 1e4
samp <- numeric(length = N)
counts <- numeric(length = N)
for(i in 1:N)
{
  rep <- gamma_ar(alpha = 4, beta = 3, lambda = 4/3) # at optimal value of lambda
  samp[i] <- rep[1]
  counts[i] <- rep[2]
}
mean(counts)

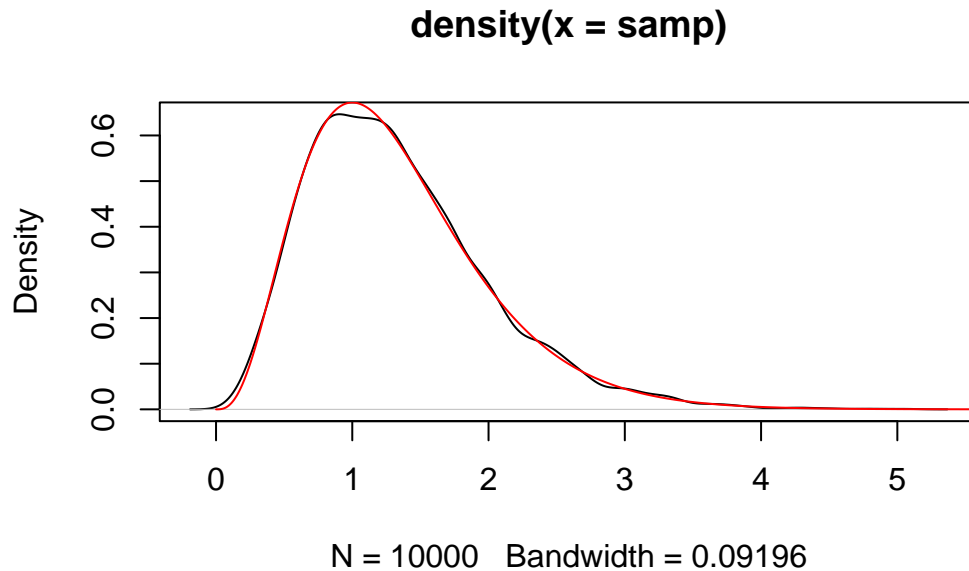
```

```
[1] 2.8656
```

```

plot(density(samp), ylab = "Density")
x <- seq(0, 6, length = 1e3)
lines(x, dgamma(x, 4, 3), col = "red")

```




The mean counts is roughly 2.9 which matches the theoretical optimal, and the density function of the sampled values matches the true density. Now we repeat for  $\lambda = 2$  which should be less efficient (c will be larger). But notice how the density matches similarly.

```
# For lambda = 2 now
N <- 1e4
samp <- numeric(length = N)
counts <- numeric(length = N)
for(i in 1:N)
{
  rep <- gamma_ar(alpha = 4, beta = 3, lambda = 2) # at value 2
  samp[i] <- rep[1]
  counts[i] <- rep[2]
}
mean(counts)
```

```
[1] 9.1226
```

```
plot(density(samp), ylab = "Density")
x <- seq(0, 6, length = 1e3)
lines(x, dgamma(x, 4, 3), col = "red")
```



3. For a  $N(0, 1)$  target, consider a Cauchy proposal with scale parameter  $\sigma$ , where the pdf of such a proposal is

$$g(x) = \frac{1}{\pi\sigma} \frac{1}{(1 + (x/\sigma)^2)}.$$

Find the optimal value of  $\sigma$ , and implement the AR algorithm for this value.

Before we do any implementation, we will have to do some theory. First, note that the target density is

$$f(x) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2}.$$

For this target and this proposal, we get

$$\begin{aligned} r(x) &:= \frac{f(x)}{g(x)} = \frac{\pi\sigma}{\sqrt{2\pi}} \left(1 + \frac{x^2}{\sigma^2}\right) e^{-x^2/2} \\ \log r(x) &= \log \frac{\pi\sigma}{\sqrt{2\pi}} + \log \left(1 + \frac{x^2}{\sigma^2}\right) - \frac{x^2}{2} \\ \Rightarrow \frac{d \log r(x)}{dx} &= \frac{2x}{\sigma^2 + x^2} - x \stackrel{\text{set}}{=} 0 \\ \Rightarrow x &= 0, x = \pm \sqrt{2 - \sigma^2} \text{ (for } \sigma^2 < 2) \end{aligned}$$

By checking the second derivatives, we can see that for  $\sigma^2 < 2$ , the maxima occurs at  $x = \pm \sqrt{2 - \sigma^2}$ , and for  $\sigma^2 > 2$ , the maxima occurs at  $x = 0$ . Thus, we obtain that the bound is

$$c(\sigma) = \sup r(x) = \max \left\{ r(0), r(\sqrt{2 - \sigma^2}) \right\} = \max \left\{ \frac{\pi\sigma}{\sqrt{2\pi}}, \frac{\pi\sigma}{\sqrt{2\pi}} \left(1 + \frac{2 - \sigma^2}{\sigma^2}\right) e^{-(2 - \sigma^2)/2} \right\}$$

In order to find the optimal value of  $\sigma^2$ , we will have to minimize  $c(\sigma)$ . We will find the minimum in two different cases:

**Case 1:** for  $\sigma^2 \geq 2$ ,  $c(\sigma)$  is an increasing function of  $\sigma$ , and thus the minimum occurs as  $\sigma^2 = 2$ , which yields the best  $c$  in this case to be  $c(\sqrt{2}) = \sqrt{\pi}$ .

**Case 2:** for  $\sigma^2 < 2$ , the minimum occurs at  $\sigma = 1$  (you can show), at which point  $c(1) = \sqrt{2\pi/e}$ .

Since  $\sqrt{\pi} > \sqrt{2\pi/e}$ , we conclude that the optimal Cauchy proposal is with  $\sigma = 1$  (the standard Cauchy)!

```
#####
## Accept-reject for
## N((2,.1))0,1) distribution
## Using Cauchy(0, sigma) proposal
#####
ARnormCauchy <- function(prop.scale = 1)
{
  if(prop.scale^2 >= 2) c <- pi*prop.scale/(sqrt(2*pi))
  if(prop.scale^2 < 2)
  {
    c <- dnorm(sqrt(2 - prop.scale^2))/dcauchy(sqrt(2 - prop.scale^2), scale = prop.scale)
  }

  accept <- 0
  counter <- 0 # count the number of loop
  while(accept == 0)
  {
    counter <- counter + 1

    # Inverse transform to draw from proposal
    prop.U <- runif(1)
    # generating from proposal -- using location scale trick
    prop <- prop.scale*tan(pi*(prop.U - .5))

    # log ratio
    log.ratio <- dnorm(prop, log = TRUE) - dcauchy(prop, scale = prop.scale, log = TRUE)
    log.ratio <- log.ratio - log(c)

    # AR step
    if(log(runif(1)) <= log.ratio)
    {
      accept <- 1
      return(c(prop, counter))
    }
  }
}
```



```

    }
  }

  ### Obtaining  $10^4$  samples from Normal() distribution
  N <- 1e4
  samp <- numeric(length = N)
  counts <- numeric(length = N)
  for(i in 1:N)
  {
    rep <- ARnormCauchy(prop.scale = 1)
    samp[i] <- rep[1]
    counts[i] <- rep[2]
  }

  # mean counts
  mean(counts)

```


```
[1] 1.5201
```

```

# Make a plot of the estimated density from the samples
# versus the true density
x <- seq(-5, 5, length = 5000)
plot(density(samp), main = "Estimated density from 1e4 samples")
lines(x, dnorm(x), col = "red", lty = 2) ## Complete this
legend("topleft", lty = 1:2, col = c("black", "red"), legend = c("AR", "truth"))

```

### Estimated density from 1e4 samples



N = 10000 Bandwidth = 0.1451

```

## Checking for other values to see if it is optimum
rowMeans(replicate(1e3, ARnormCauchy(prop.scale = 2)))[2]

[1] 2.511

rowMeans(replicate(1e3, ARnormCauchy(prop.scale = .8)))[2]

[1] 1.589

rowMeans(replicate(1e3, ARnormCauchy(prop.scale = 10)))[2]

[1] 12.09

rowMeans(replicate(1e3, ARnormCauchy(prop.scale = .5)))[2]

[1] 2.035

```

Empirically as well, we see that  $\sigma = 1$  gives the optimal proposal.