

MTH210a: Lab 3 Solutions

1. The file `BetaAR.R` contains partial code to implement an AR algorithm for a `Beta(4,3)` target. Complete the code and analyse the results.

Below is the complete code, along with the code for plots.

```
#####  
## Accept-reject for  
## Beta(4,3) distribution  
## Using U(0,1) proposal  
#####  
set.seed(1)  
beta_ar <- function()  
{  
  c <- 60 * (3/5)^3 * (2/5)^2  
  accept <- 0  
  counter <- 0 # count the number of loop  
  while(accept == 0)  
  {  
    counter <- counter + 1  
  
    prop <- runif(1)  
    ratio <- dbeta(prop, 4, 3)/c  
  
    U <- runif(1)  
  
    if(U <= ratio)  
    {  
      accept <- 1  
      return(c(prop, counter))  
    }  
  }  
}  
  
### Obtaining 10^4 samples from Beta() distribution
```


```

N <- 1e4
samp <- numeric(length = N)
counts <- numeric(length = N)
for(i in 1:N)
{
  rep <- beta_ar() ## fill in
  samp[i] <- rep[1] ## fill in
  counts[i] <- rep[2] ## fill in
}

# Make a plot of the estimated density from the samples
# versus the true density
x <- seq(0, 1, length = 500)
plot(density(samp), main = "Estimated density from 1e4 samples")
lines(x, dbeta(x, 4, 3), col = "red", lty = 2) ## Complete this
legend("topleft", lty = 1:2, col = c("black", "red"), legend = c("AR", "truth"))

```

Estimated density from 1e4 samples



N = 10000 Bandwidth = 0.02489

```

# This is c
(c <- 60 * (3/5)^3 * (2/5)^2)

```

```
[1] 2.0736
```

```

# This is the mean number of loops required
mean(counts)

```

```
[1] 2.0936
```

```
#They should be almost the same!
```

2. Write R code for Problem 7 in Exercises from Section 4 of the notes.

Let X be an $\text{Exp}(1)$. Provide an efficient algorithm for simulating a random variable whose distribution is the conditional distribution of X given that $X < 0.05$. That is, its density function is

$$f(x) = \frac{e^{-x}}{1 - e^{-0.05}} \quad 0 < x < 0.05.$$

Using R generate 1000 such random variables and use them to estimate $E[X \mid X < 0.05]$.

First, we will do the theory for this. Note that the target density is the truncated exponential(1) truncated to be between 0 and 0.05. Just like the previous truncation examples, an AR is easy, if we use an Exponential proposal. I will not show the math for this; please do this by your self. We will get finally for $Y \sim \text{Exp}(1)$

$$\frac{f(y)}{cg(y)} = I(0 \leq y \leq .05)$$

After we get samples from the truncated exponential, we need to return the mean of this truncated exponential. We can estimate the population mean with the sample mean: so finally when we get X_1, X_2, \dots, X_n from Truncated Exponential, we will then estimate $E[X \mid X < 0.05]$ with

$$\frac{1}{n} \sum_{t=1}^n X_t$$

```
#### sample from truncated exp
truncExp <- function()
{
  accept <- 0
  count <- 0
  # inverse transform
  # to sample from exp
  while(!accept)
  {
    count <- count + 1
    U <- runif(1)
    expo <- -log(U)
    if(expo <= .05)
    {
      accept <- 1
      return(c(expo, count))
    }
  }
}
```

```

    }
  }
}

## Obtaining multiple samples
N <- 1e4
samples <- numeric(length = N)
try <- numeric(length = N)
for(i in 1:N)
{
  rep <- truncExp()
  samples[i] <- rep[1]
  try[i] <- rep[2]
}
mean(samples) # answer

```

```
[1] 0.02476953
```

3. The file `circleAR.R` contains partial code to implement the accept-reject sampler to draw from the uniform distribution over the circle. Complete the code.

```

#####
## Accept-reject for obtaining
## sample uniformly from a standard circle
## using a box as a proposal
#####
set.seed(1)
circle_ar <- function()
{
  accept <- 0
  counter <- 0 # count the number of loop
  while(accept == 0)
  {
    counter <- counter + 1

    prop.temp <- runif(2) # from U(0,1)
    prop <- -1 + 2*prop.temp # from U(-1,1)

    if(prop[1]^2 + prop[2]^2 <= 1) # fill condition
    {
      accept <- 1
      return(c(prop, counter))
    }
  }
}

```

```

    }
  }

  # Simulation 10^4 samples from circle
  N <- 1e4
  samp <- matrix(0, ncol = 2, nrow = N)
  counts <- numeric(length = N)
  for(i in 1:N)
  {
    foo <- circle_ar() # I use foo as a dummy name
    samp[i,] <- foo[1:2]
    counts[i] <- foo[3]
  }

  4/pi

```

```
[1] 1.27324
```

```

# [1] 1.27324
mean(counts) # should be very close

```


```
[1] 1.2783
```

```

# Plotting the obtained samples
# no particular part of the circle is favored more
# than any other part.
plot(samp[,1], samp[,2], xlab = "x", ylab = "y",
     main = "Uniform samples from a circle", asp = 1)

```

Uniform samples from a circle



4. Taking inspiration from `circleAR.R`, implement Problem 16 from Section 4 Exercises of the notes.

- a. Implement an accept-reject sampler to sample uniformly from the circle $\{x^2 + y^2 \leq 1\}$ and obtain 10000 samples and estimate the probability of acceptance. Does it approximately equal $\pi/4$?

We have already done this.

- b. Now consider sampling uniformly from a p -dimensional sphere (a circle is $p = 2$). Consider a p -vector $\mathbf{x} = (x_1, x_2, \dots, x_p)$ and let $\|\cdot\|$ denote the Euclidean norm. The pdf of this distribution is

$$f(\mathbf{x}) = \frac{\Gamma(\frac{p}{2} + 1)}{\pi^{p/2}} I\{\|\mathbf{x}\| \leq 1\}.$$

Use a uniform p -dimensional hypercube to sample uniformly from this sphere. Implement this for $p = 3, 4, 5$, and 6 . What happens as p increases?

To code this question, we will first have to do some theory to ensure that c remains finite and to understand what the value of c will be. We consider a p dimensional box as the proposal, centered at the origin. The pdf of the uniform distribution over this box is

$$g(\mathbf{x}) = \frac{1}{2^p} I(-1 \leq x_i \leq 1, i = 1, \dots, p).$$

For this, we can find c since

$$\sup_{\mathbf{x}} \frac{f(\mathbf{x})}{g(\mathbf{x})} = \frac{\Gamma(\frac{p}{2} + 1) 2^p}{\pi^{p/2}} I\{\|\mathbf{x}\| \leq 1\} \leq \frac{\Gamma(\frac{p}{2} + 1) 2^p}{\pi^{p/2}}.$$

The above value of c increases rapidly as a function of p

```
c_sphere <- function(p)
{
  gamma(p/2 + 1)* 2^p/ (pi^(p/2))
}
c_sphere(c(2:6, 10, 30))
```

```
[1] 1.273240e+00 1.909859e+00 3.242278e+00 6.079271e+00 1.238459e+01
```

```
[6] 4.015428e+02 4.899496e+13
```

The value of c increases rapidly with p . So we can see that the algorithm will slow down incredibly in higher dimensions.

```
#####
## Accept-reject for obtaining
## sample uniformly from a sphere
## using a box as a proposal
#####
sphere_ar <- function(p = 3)
{
  accept <- 0
  counter <- 0 # count the number of loop
  while(accept == 0)
  {
    counter <- counter + 1

    prop.temp <- runif(p) # from U(0,1)
    prop <- -1 + 2*prop.temp # from U(-1,1)

    if(sum(prop^2) <= 1) # fill condition
    {
      accept <- 1
      return(c(prop, counter))
    }
  }
}

# Simulation 10^3 samples from circle
N <- 1e3
p <- 4
samp <- matrix(0, ncol = p+1, nrow = N)
counts <- numeric(length = N)
for(i in 1:N)
```

```
{
  foo <- sphere_ar(p = p) # I use foo as a dummy name
  samp[i, 1:p] <- foo[1:p]
  counts[i] <- foo[p+1]
}
```

```
c_sphere(p = 4)
```

```
[1] 3.242278
```

```
mean(counts) # should be very close
```

```
[1] 3.136
```

5. Will share solutions later – once taught in class.
6. Will share solutions later - once taught in class
7. **Suppose $Y = \sum_{i=1}^5 X_i$ where $X_i \sim \text{Weibull}(\alpha_i, \lambda)$. Here density of $\text{Weibull}(\alpha, \lambda)$ is**

$$f(x) = \alpha \lambda^{-\alpha} x^{\alpha-1} e^{-(x/\lambda)^\alpha}, \quad x > 0.$$

Using only $U(0, 1)$ draws, estimate $E(Y^2)$. Assume $\alpha_i = i$ and $\lambda = 5$.

Unfortunately there were a few typos in the question. The above is the corrected density. First, I need to figure out how to sample from Weibull distribution. Since inverse transform is fairly easy method, I want to try that first. Using change of variables trick, it can be shown that

$$F(x) = 1 - e^{-(x/\lambda)^\alpha}.$$

Inverting this function, I obtain that

$$F^{-1}(u) = \lambda [-\log(1 - u)]^{1/\alpha}.$$

This means, we can sample from Weibull easily. So in order to estimate $E(Y^2)$, I note that if I can obtain $Y_1, Y_2, \dots, Y_n \stackrel{iid}{\sim}$ Distribution of Y , then I can estimate this expectation with:

$$\frac{1}{n} \sum_{t=1}^n Y_t^2$$

Simulating from the distribution of Y is possible by sampling Weibulls and adding them up as the formula indicates. Below, the function `distY` obtains one draw from Y given a vector of α and λ .


```
#####
# Sample from dist of Y
#####
distY <- function(alpha, lambda)
{
  l <- length(alpha)
  Wi <- numeric(length = l)
  for(i in 1:l)
  {
    U <- runif(1)
    Wi[i] <- lambda*(-log(1-U))^(1/alpha[i])
  }
  return(sum(Wi))
}
```

Now, I will call this function $n = 1e3$ times to estimate $E(Y^2)$ from a sample average of these 1000 Y 's.

```
### Estimate expectation with average
samples <- replicate(1e3, distY(alpha = 1:5, lambda = 5))

## Final answer
mean(samples^2)
```

```
[1] 586.3119
```

```
## Just for information, here is a
## hist of samples of Y
hist(samples, main = "Histogram of samples from Y")
```

Histogram of samples from Y

